
Harmony-Driven Theory Discovery in Knowledge Graphs

via LLM-Guided Island Search

Anonymous Author(s)

Affiliation

Address

email

Abstract

Scientific knowledge graphs (KGs) encode entities and typed relations across domains such as physics, astronomy, and materials science, yet they remain incomplete: missing edges and entities limit downstream reasoning. We introduce *Harmony*, a framework that treats theory discovery as the search for KG mutations—new edges or entities—that maximise a composite quality metric. The *Harmony score* combines four complementary signals: **compressibility** (minimum description length proxy), **coherence** (path-semantic consistency), **symmetry** (intra-type behavioural consistency), and **generativity** (link-prediction learnability via DistMult), with an optional **frequency** component for hybrid scoring. An LLM proposer generates candidate theory-level propositions, which are validated, scored, and archived in a MAP-Elites quality-diversity grid across an island-model search topology. We evaluate on seven KG domains (five hand-curated, two Wikidata-sourced) with 10-seed cross-validation and four external KGE baselines (DistMult, TransE, RotatE, ComplEx). Factor decomposition isolates the contributions of the LLM proposer, Harmony scoring, and quality-diversity archive. A regime characterisation analysis reveals that Harmony adds complementary value to frequency-based approaches in denser, multi-type KGs: on Wikidata Materials, Harmony-guided proposals show the largest directional improvement over DistMult-alone on both Hits@10 (0.34 vs. 0.32) and MRR (0.12 vs. 0.11), though differences do not reach statistical significance at $\alpha = 0.05$. Backtesting against held-out edges shows that proposals do not reproduce known edges, serving as a negative control consistent with theory-level rather than interpolation behaviour; a typed-neighbourhood soft-match diagnostic suggests proposals tend toward structurally adjacent knowledge. We frame Harmony as a methodology for theory-shaped proposal generation; assessing whether produced propositions constitute novel scientific discoveries requires human expert evaluation, which we defer to future work.

1 Introduction

Knowledge graphs (KGs) organise scientific knowledge as typed, directed multigraphs: entities represent concepts (e.g. *photon*, *eigenvalue*, *graphene*) and edges encode semantic relations such as *derives*, *explains*, or *contradicts* [5]. Despite decades of curation, scientific KGs remain structurally incomplete—missing edges that encode latent theoretical connections and missing entities that represent undiscovered concepts.

Knowledge graph completion (KGC) methods—TransE [2], DistMult [14], RotatE [12]—learn low-dimensional embeddings and predict missing links. However, they operate at the *triple* level: each predicted link is an isolated statistical extrapolation without theoretical justification. They do not

36 produce *theory-level propositions* that articulate *why* a relation should hold, what it implies, or how
37 it could be falsified.

38 We address this gap with **Harmony**, a framework for automated theory discovery in scientific KGs.
39 The key idea is a composite quality metric—the *Harmony score*—that captures four desiderata of a
40 well-structured knowledge graph:

- 41 1. **Compressibility**: the KG’s edge-type distribution and spanning structure admit a short
42 description (MDL proxy).
- 43 2. **Coherence**: closed paths exhibit consistent edge-type semantics and contradictions are
44 sparse.
- 45 3. **Symmetry**: entities of the same type use edge types in similar proportions (low Jensen–
46 Shannon divergence).
- 47 4. **Generativity**: a shallow DistMult model can recover masked edges, indicating learnable
48 relational patterns.

49 An LLM proposes candidate mutations—adding edges or entities—each accompanied by a claim,
50 justification, and falsification condition. Proposals are scored by Harmony gain and archived in a
51 MAP-Elites [9] quality-diversity grid across an island-model [13] search topology (Section 3.4).

52 Contributions.

- 53 1. A four-component **Harmony metric** for scoring KG quality that is domain-agnostic,
54 bounded in $[0, 1]$, and decomposes into interpretable sub-scores (Section 3.2).
- 55 2. A **proposal schema** that elevates KG mutations from bare triples to falsifiable theory-level
56 claims (Section 3.3).
- 57 3. An **island-model LLM search loop** with MAP-Elites archiving and stagnation-triggered
58 constrained prompting (Section 3.4).
- 59 4. Empirical evaluation on **seven KG domains** (five hand-curated: linear algebra, periodic ta-
60 ble, astronomy, physics, materials science; two Wikidata-sourced: Wikidata Physics, Wiki-
61 data Materials). Harmony-guided proposals show directional improvements over DistMult-
62 alone on three of five discovery domains (Materials, Wikidata Physics, Wikidata Materials),
63 none reaching statistical significance at $\alpha = 0.05$ and with small effect sizes ($|\delta_C| \leq 0.12$);
64 the framework’s contribution is the factor decomposition isolating the MAP-Elites archive
65 as the dominant driver of structural gain (Section 5).

66 2 Related Work

67 **Knowledge graph completion.** Embedding-based methods—TransE [2], DistMult [14], RotatE
68 [12], among others [5]—project entities and relations into low-dimensional vector spaces and pro-
69 duce ranked link predictions without theoretical justification. Our work uses DistMult as the genera-
70 tivity *component* within a broader metric, and additionally generates natural-language propositions
71 explaining each mutation.

72 **Automated scientific discovery.** FunSearch [10] uses LLMs to discover mathematical construc-
73 tions by evolving Python programs. PySR [4] performs symbolic regression via genetic program-
74 ming [7], discovering closed-form expressions from numerical data. The survey by Makke and
75 Chawla [8] covers the broader symbolic regression landscape. These systems discover *formulas*
76 over numerical features; Harmony discovers *relational propositions* over typed knowledge graphs,
77 a structurally different search space.

78 **Quality-diversity search.** MAP-Elites [9] maintains a grid of solutions indexed by behavioural
79 descriptors. We adopt MAP-Elites with a two-dimensional descriptor (simplicity, Harmony gain)
80 combined with an island-model [13] topology (Section 3.4).

81 **LLM-guided reasoning over KGs.** Recent work integrates LLMs with structured knowledge
82 graphs in several ways. KAPING [1] augments LLM prompts with retrieved KG triples for zero-shot
83 question answering. Think-on-Graph [11] performs multi-hop reasoning by iteratively traversing
84 KG neighbours guided by LLM chain-of-thought. StructGPT [6] provides a general interface for
85 LLMs to query and reason over structured data including KGs. These systems use KGs as *context*

for LLM reasoning; our approach inverts the role: the LLM is a *proposer* that generates structured mutations (new edges and entities) with accompanying justifications, and a deterministic Harmony metric—not LLM self-evaluation—scores and selects proposals.

3 Method

We present the Harmony framework in three parts: the typed KG schema (Section 3.1) and Harmony metric (Section 3.2), the proposal schema and validation (Section 3.3), and the island-model search loop (Section 3.4).

3.1 Typed Knowledge Graph Schema

A knowledge graph $G = (V, E)$ consists of entities V and typed directed edges E . Each entity $v \in V$ has an `entity_type` label (e.g. *concept*, *element*, *celestial_object*) and a property bag. Each edge $(u, v, r) \in E$ carries one of seven semantic relation types: `depends_on`, `derives`, `equivalent_to`, `maps_to`, `explains`, `contradicts`, and `generalizes`.

The seven relation types cover dependency, derivation, equivalence, correspondence, explanation, contradiction, and hierarchy—a minimal set inspired by morphism classes in category theory (Appendix O).

3.2 Harmony Metric

The Harmony score combines four core signals, each normalised to $[0, 1]$, with an optional frequency component:

$\mathcal{H}(G) = \alpha \cdot \text{Compress}(G) + \beta \cdot \text{Cohere}(G) + \gamma \cdot \text{Symm}(G) + \delta \cdot \text{Gener}(G) + \varepsilon \cdot \text{Freq}(G)$, (1)
where $\alpha, \beta, \gamma, \delta, \varepsilon \geq 0$ are normalised internally so that $\alpha + \beta + \gamma + \delta + \varepsilon = 1$. Default weights are uniform over the four core components ($\alpha = \beta = \gamma = \delta = 0.25, \varepsilon = 0$); the frequency component $\text{Freq}(G)$ is the Hits@ K of the frequency heuristic (predicting the most common edge type between entity pairs) and is included only when $\varepsilon > 0$ for hybrid analysis (Appendix I).

Compressibility. An MDL proxy combining edge-type entropy and spanning structure:

$$\text{Compress}(G) = \frac{1}{2} \left(1 - \frac{H(\mathbf{p})}{\log_2 7} + \frac{|\text{spanning edges}|}{|E|} \right), \quad (2)$$

where $H(\mathbf{p})$ is the Shannon entropy of the edge-type distribution and the spanning fraction counts BFS spanning-tree edges over an undirected view of G .

Coherence. Path-semantic consistency via triangle type-agreement and contradiction rate:

$$\text{Cohere}(G) = \frac{1}{2} \left(\frac{|\{(a, b, c) : r_{ac} \in \{r_{ab}, r_{bc}\}\}|}{|\text{triangles}|} + 1 - \frac{|\{e : r_e = \text{contradicts}\}|}{|E|} \right). \quad (3)$$

Symmetry. Intra-type behavioural consistency via Jensen–Shannon divergence. For each entity type τ , let $\mathbf{q}_v \in \Delta^6$ be entity v ’s outgoing edge-type distribution:

$$\text{Symm}(G) = \frac{\sum_{\tau} n_{\tau} \cdot \left(1 - \frac{1}{n_{\tau}} \sum_{v \in \tau} \text{JS}(\mathbf{q}_v, \bar{\mathbf{q}}_{\tau}) \right)}{\sum_{\tau} n_{\tau}}. \quad (4)$$

Generativity. Link-prediction learnability via DistMult [14] on 20% masked edges:

$$\text{Gener}(G) = \text{Hits@}K(\text{DistMult}, G_{\text{mask}}), \quad (5)$$

with $d=50$ embeddings, 100 training epochs, and $K=10$.

Proposal value function. Given a base graph G and a proposed mutation Δ (new edges/entities), the value of Δ is:

$$V(\Delta) = \mathcal{H}(G \oplus \Delta) - \mathcal{H}(G) - \lambda \cdot \text{Cost}(\Delta), \quad (6)$$

where $G \oplus \Delta$ denotes the graph after applying Δ , and $\text{Cost}(\Delta) = 1/\max(|E|, 1)$ for edge mutations, $1/\max(|V|, 1)$ for entity mutations. The penalty weight $\lambda = 0.1$ discourages trivially large proposals. The Harmony score is bounded in $[0, 1]$, decomposes into independently computable components, and exhibits empirical directional monotonicity (Appendix P).

3.3 Proposal Schema and Validation

Each proposal is a structured record containing:

- **Mutation type:** ADD_EDGE, REMOVE_EDGE, ADD_ENTITY, or REMOVE_ENTITY.
- **Claim:** a one-sentence theoretical statement (e.g. “Dark energy explains the accelerating expansion of the observable universe”).
- **Justification:** reasoning supporting the claim.
- **Falsification condition:** what evidence would disprove the claim.
- **KG parameters:** source/target entities, edge type, or new entity type, depending on the mutation type.

A deterministic validator enforces three rules: (i) text fields must be ≥ 10 characters, (ii) type-specific parameters must be present (e.g. ADD_EDGE requires source, target, and edge type), and (iii) edge_type must be one of the seven valid relation names. Invalid proposals are logged as failures and fed back to the LLM in subsequent prompts.

3.4 Island-Model Search with MAP-Elites

Island topology. Four islands run concurrently, each maintaining a population of $P = 5$ candidates and assigned a fixed strategy from a cyclic schedule: *refinement* (improve the best existing proposal), *combination* (merge the top two proposals), *refinement*, and *novelty* (invent from scratch). Each island uses a distinct LLM temperature: $\{0.3, 0.3, 0.8, 1.2\}$ to further diversify exploration.

MAP-Elites archive. A shared 5×5 MAP-Elites grid [9] indexes proposals by two behavioural descriptors: *simplicity* (inverse structural cost) and *Harmony gain* ($\mathcal{H}(G \oplus \Delta) - \mathcal{H}(G)$). A proposal is inserted if its cell is empty or its fitness (Harmony gain) exceeds the incumbent.

Stagnation recovery. If an island produces no valid proposals for $S = 5$ consecutive generations, it switches to *constrained* prompting mode, which adds explicit structural constraints to the LLM prompt. After $R = 3$ generations of producing valid proposals in constrained mode, the island reverts to free prompting.

Migration. Every $M = 10$ generations, the best proposal from each island migrates to the next island in a ring topology (island $i \rightarrow \text{island } (i + 1) \bmod 4$), where it is prepended to the destination island’s prompt-context buffer for the next generation.

Generation loop. Algorithm 1 summarises a single generation (Figure 1). The loop runs for $T_{\max} = 20$ generations per experiment, checkpointing state after each generation to enable resumption.

4 Experiments

4.1 Knowledge Graph Domains

We evaluate on seven KGs: five hand-curated (linear algebra, periodic table, astronomy, physics, materials science) and two Wikidata-sourced (Wikidata Physics, Wikidata Materials). The hand-curated KGs range from 41–153 entities; Wikidata KGs have 252–253 entities and 800+ edges. All use the shared seven-relation type vocabulary (Section 3.1). Full statistics are in Table 2. The first two domains serve as calibration targets; the remaining five are discovery targets.

4.2 Dataset Splitting

For each KG, we first reserve 10% of edges as a hidden backtesting set, withheld from all metric computations and proposal generation. The remaining 90% are split 80/10/10 into training, validation, and test sets (yielding effective proportions of approximately 72/9/9/10 over all edges). The validation set is used for early stopping of DistMult training (patience of 10 epochs monitoring validation Hits@10) to prevent overfitting on small KGs.

Algorithm 1 Harmony search — one generation

Require: Base KG G , islands $\{I_1, \dots, I_4\}$, archive $\mathcal{A}$

```
1: for each island  $I_k$  do
2:    $\sigma_k \leftarrow \text{STRATEGY}(k)$  {refinement / combination / novelty}
3:   prompt  $\leftarrow \text{BUILDPROMPT}(G, \sigma_k, \text{top}(I_k), \text{failures}(I_k))$ 
4:    $\hat{p} \leftarrow \text{LLM}(\text{prompt}, \text{temp}_k)$ 
5:   if  $\text{VALIDATE}(\hat{p})$  then
6:      $\Delta \leftarrow \text{apply } \hat{p} \text{ to } G$ 
7:      $v \leftarrow V(\Delta)$  {Eq. 6}
8:      $\text{TRYINSERT}(\mathcal{A}, \hat{p}, v, \text{descriptor}(\hat{p}))$  {descriptor = (simplicity, gain)}
9:     Update  $I_k$  population
10:  else
11:    Log failure; feed back to next prompt
12:  end if
13: end for
14: if generation mod  $M = 0$  then
15:    $\text{MIGRATE TOP1 TO CONTEXT}(I_1, \dots, I_4)$  {ring topology}
16: end if
```




Figure 1: Overview of the Harmony island-model search framework. Four islands run concurrently, each proposing KG mutations via an LLM following a cyclic strategy schedule (Refine $\rightarrow$ Combine $\rightarrow$ Refine $\rightarrow$ Novel). Proposals are scored by the Harmony metric and archived in a MAP-Elites quality-diversity grid (5×5 , axes = simplicity bin $\times$ Harmony-gain bin). Periodic migration (every $M = 10$ generations) transfers each island’s best proposal to the next island in a ring topology ($i \rightarrow (i+1) \bmod 4$).

4.3 Baselines

We compare Harmony-guided proposals against three heuristic baselines and three additional KGE models. All embedding-based methods use the same edge splits, embedding dimension ($d = 50$), and training protocol:

1. **Random:** random entity pairs with random relation types.
2. **Frequency:** most frequent relation type between most-connected entity pairs.
3. **DistMult-alone:** DistMult’s own top-ranked predictions without Harmony scoring.
4. **TransE, RotatE, ComplEx:** translational, rotation-based, and complex-valued KGE models respectively.

We additionally report **Harmony-3** ($\delta = 0$, no generativity) to address evaluation circularity; it matches DistMult-alone on all domains, confirming that generativity drives any observed difference. We also run three ablated configurations (LLM-only, Harmony-only, No-QD) to isolate component contributions (Appendix S).




Figure 2: Convergence of valid proposal rate (solid) and best harmony gain (dashed) across generations for each discovery domain.

4.4 Evaluation Protocol

Quantitative metrics. We report Hits@10 and Mean Reciprocal Rank ($MRR = \frac{1}{|Q|} \sum_i 1/\text{rank}_i$) on the test split after applying top proposals from the MAP-Elites archive to the base KG. All main results (Table 1) are averaged over 10 random seeds. LLM proposals are generated by `gpt-oss:20b`, a dense 20B-parameter open-source model locally served via Ollama. Factor decomposition experiments use `mlx-community/Qwen3.5-35B-A3B-4bit` via `mlx_lm` at ~ 46 tokens/s. All inference runs on a single Apple M2 Pro CPU.

Additional protocols. The calibration gate protocol (Appendix Q), backtesting protocol (Appendix T), and archive diversity protocol (Appendix R) provide complementary validation beyond link prediction metrics.

5 Results

Table 1 compares link prediction metrics across five discovery domains after applying top proposals from the MAP-Elites archive to the base KG.

Link prediction. Harmony-guided proposals show directional improvements over DistMult-alone on Hits@10 in three of five domains—Wikidata Materials (0.34 vs. 0.32), Materials (0.31 vs. 0.29), and Wikidata Physics (0.26 vs. 0.25)—but none reach statistical significance ($p = 0.38\text{--}1.00$, $|\delta_C| \leq 0.12$; Section H). TransE achieves the highest Hits@10 on three domains; frequency leads on two. On Physics, Harmony underperforms DistMult-alone (0.32 vs. 0.37), due to the LLM’s tendency to propose edges between hub entities (Section 6). Variance is substantially lower on the denser Wikidata KGs (std $\approx 0.01\text{--}0.05$) than on the smaller hand-curated domains (std $\approx 0.04\text{--}0.23$).

Convergence and archive coverage. Across all domains, the valid proposal rate reaches ≥ 0.50 by generation 10 (Figure 2). The MAP-Elites archive achieves 40–60% cell coverage, indicating diverse proposals spanning multiple simplicity–gain trade-offs (Appendix R).

Factor decomposition. Three ablated configurations isolate component contributions (Appendix S). Key findings: (i) the MAP-Elites archive is the dominant driver of structural gain—Harmony-only (random proposer) outperforms No-QD (greedy archive) on all five domains; (ii) LLM-only (no Harmony scoring) saturates trivially, confirming that Harmony scoring provides essential selection pressure; (iii) the LLM proposer’s primary value is *semantic*—producing justified, falsifiable claims—rather than maximising raw Harmony gain.

Backtesting. Backtesting against 10% held-out edges yields zero exact matches across all domains, confirming proposals behave as theory-level mutations rather than triple interpolation. Soft-match rates reach 10–25% on structurally adjacent neighbourhoods (Appendix T).

Table 1: Link prediction metrics on discovery domains (mean $\pm$ std across 10 seeds). Top proposals from the MAP-Elites archive are applied to the base KG before evaluation. Best Hits@10 per domain in **bold**; best MRR underlined. p -values from paired bootstrap CI (2000 resamples) comparing Harmony vs. DistMult-alone; δ_C is Cliff’s delta effect size.[†]

Domain	Method	Hits@10	MRR	p	δ_C
Astronomy	Random	0.27 $\pm$ 0.16	0.12 $\pm$ 0.10	1.00	+0.00
	Frequency	0.39 $\pm$ 0.12	<u>0.16</u> $\pm$ 0.08		
	TransE	0.33 $\pm$ 0.22	0.11 $\pm$ 0.05		
	RotatE	0.16 $\pm$ 0.16	0.09 $\pm$ 0.07		
	ComplEx	0.33 $\pm$ 0.21	0.11 $\pm$ 0.07		
	DistMult-alone	0.24 $\pm$ 0.17	0.10 $\pm$ 0.04		
	Harmony (ours)	0.24 $\pm$ 0.17	0.10 $\pm$ 0.04		
Physics	Random	0.29 $\pm$ 0.13	0.10 $\pm$ 0.07	0.62	−0.06
	Frequency	0.46 $\pm$ 0.12	<u>0.25</u> $\pm$ 0.08		
	TransE	0.50 $\pm$ 0.11	0.20 $\pm$ 0.06		
	RotatE	0.29 $\pm$ 0.20	0.11 $\pm$ 0.06		
	ComplEx	0.38 $\pm$ 0.16	0.22 $\pm$ 0.11		
	DistMult-alone	0.37 $\pm$ 0.14	0.16 $\pm$ 0.07		
	Harmony (ours)	0.32 $\pm$ 0.23	0.13 $\pm$ 0.09		
Materials	Random	0.17 $\pm$ 0.12	0.11 $\pm$ 0.06	0.72	+0.06
	Frequency	0.36 $\pm$ 0.18	0.15 $\pm$ 0.08		
	TransE	0.42 $\pm$ 0.17	<u>0.17</u> $\pm$ 0.08		
	RotatE	0.26 $\pm$ 0.10	0.09 $\pm$ 0.04		
	ComplEx	0.32 $\pm$ 0.10	0.13 $\pm$ 0.04		
	DistMult-alone	0.29 $\pm$ 0.14	0.15 $\pm$ 0.09		
	Harmony (ours)	0.31 $\pm$ 0.14	0.13 $\pm$ 0.05		
Wikidata Physics	Random	0.05 $\pm$ 0.01	0.02 $\pm$ 0.01	0.89	−0.06
	Frequency	0.29 $\pm$ 0.02	<u>0.14</u> $\pm$ 0.02		
	TransE	0.30 $\pm$ 0.03	0.11 $\pm$ 0.01		
	RotatE	0.18 $\pm$ 0.02	0.07 $\pm$ 0.01		
	ComplEx	0.25 $\pm$ 0.03	0.09 $\pm$ 0.01		
	DistMult-alone	0.25 $\pm$ 0.02	0.10 $\pm$ 0.01		
	Harmony (ours)	0.26 $\pm$ 0.04	0.09 $\pm$ 0.02		
Wikidata Materials	Random	0.03 $\pm$ 0.02	0.02 $\pm$ 0.01	0.38	+0.12
	Frequency	0.39 $\pm$ 0.03	<u>0.20</u> $\pm$ 0.02		
	TransE	0.30 $\pm$ 0.04	0.10 $\pm$ 0.01		
	RotatE	0.16 $\pm$ 0.03	0.06 $\pm$ 0.01		
	ComplEx	0.29 $\pm$ 0.03	0.10 $\pm$ 0.02		
	DistMult-alone	0.32 $\pm$ 0.05	0.11 $\pm$ 0.02		
	Harmony (ours)	0.34 $\pm$ 0.04	0.12 $\pm$ 0.01		

[†]Harmony-3 ($\delta = 0$) matches DistMult-alone on all domains, confirming that generativity drives any observed difference; full rows in Appendix B.

211 **Automated quality assessment.** Automated scoring (Gemini 2.5 Pro, LLM-as-judge) rates top-
212 5 proposals at 4.5/5 falsifiability and 4.6/5 clarity, with moderate novelty (2.8/5); full results in
213 Appendix V. Human expert evaluation with measured inter-annotator agreement is planned.

214 **Calibration and regime.** The calibration gate passed on both calibration domains (31–65% over
215 frequency; Appendix Q). Harmony’s complementary value increases with graph density and edge-
216 type diversity (Appendix U). Per-component ablation details are in Appendix B.

6 Discussion

Compressibility–generativity tension. Adding edges typically *reduces* compressibility (the BFS spanning fraction drops) while potentially *improving* generativity (more training signal for DistMult). This tension is by design: the Harmony metric rewards proposals that improve learnability without degrading structural simplicity. The value function (Eq. 6) with $\lambda > 0$ further penalises large mutations.

Frequency dominance. The frequency heuristic is highly competitive, leading on two of five domains. This is expected in small KGs with peaked edge-type distributions; we characterise the regime where Harmony adds complementary value (Appendix U).

Per-domain failure analysis. On Physics, Harmony underperforms DistMult-alone. The LLM tends to propose edges between highly connected hub entities (force, energy) rather than peripheral concepts, producing valid but redundant proposals. This suggests domain-specific prompting strategies may be needed for hub-dominated KGs.

LLM dependence and safety. The main proposer is a single model (gpt-oss:20b via Ollama); factor-decomposition runs use mlx-community/Qwen3.5-35B-A3B-4bit via mlx_lm, and the automated rubric uses Gemini 2.5 Pro (Appendix N). Ensembling across model families within the proposer role could improve diversity. To mitigate LLM-generated misinformation, proposals enter a staging layer: they are scored and archived but never automatically integrated into the base KG. Every proposal requires a falsification condition, enabling principled rejection. Total compute is ~ 10 min/domain on a single CPU (Appendix M).

Broader impacts. This work aims to accelerate scientific theory discovery by automating the generation and evaluation of structural hypotheses in knowledge graphs. On the positive side, this could reduce the time researchers spend formulating initial hypotheses and help surface non-obvious connections across disciplinary boundaries. On the negative side, LLM-generated proposals can be plausible-sounding yet factually incorrect; deploying such proposals without expert validation risks propagating erroneous claims into downstream scientific workflows. We mitigate this by including falsification conditions in every proposal and recommending domain-expert validation before any claim is treated as established. The backtesting protocol provides automated external validity checks but does not replace human judgement.

Limitations. (i) The seven-relation type vocabulary may be too coarse for highly specialised fields. (ii) The Harmony metric treats all edge types equally in compressibility and coherence; domain-specific type hierarchies could improve these signals. (iii) Observed differences between Harmony and DistMult-alone are small relative to cross-seed variance; larger-scale experiments may be needed for stronger statistical significance. (iv) The frequency heuristic remains highly competitive, leading on Hits@10 on two of five discovery domains (Astronomy and Wikidata Materials) and ranking within the top two on the remainder; this advantage is most pronounced on small KGs with peaked edge-type distributions. Harmony does not displace frequency baselines; rather, the regime characterisation (Section 5) shows that Harmony’s complementary value over DistMult-alone is most visible in denser, multi-type KGs such as Wikidata Materials, where directional improvement is observed despite frequency remaining the strongest single baseline on that domain. (v) Evaluation circularity is partially addressed by Harmony-3 and external KGE baselines, but the primary configuration still includes DistMult. (vi) Additional analysis—sparse KG challenges, proposal quality vs. validity rate, symmetry redesign rationale, and Harmony–downstream correlation—is in Appendix W.

7 Conclusion

We presented Harmony, a framework for automated theory discovery in scientific knowledge graphs combining a four-component quality metric with LLM-guided proposal generation and MAP-Elites archiving across an island-model search topology.

Evaluation across seven KG domains with 10-seed cross-validation and four external KGE baselines shows that Harmony-guided proposals achieve directional improvements over DistMult-alone on

Hits@10 in three of five domains, though differences do not reach statistical significance. The frequency heuristic remains highly competitive on small KGs. Factor decomposition identifies the MAP-Elites archive as the dominant driver of structural gain, while the LLM proposer’s primary contribution is semantic—producing justified, falsifiable theory-level claims.

Future work includes scaling to larger scientific KGs, domain-adaptive component weighting, multi-LLM ensembles across islands, and human expert evaluation with measured inter-annotator agreement.

A Dataset Statistics and Reproducibility

Table 2 provides a unified summary of all seven KG domains, including data sources and construction methods. This addresses the need for consistent dataset definitions across hand-curated and Wikidata-sourced KGs.

Table 2: Unified reproducibility table for all seven KG domains. All KGs use the shared seven-relation type vocabulary. Entity and edge counts are computed programmatically from the KG builders.

Domain	Entities	Edges	Entity types	Edge types	Source
Linear algebra	50	75	2	7	Hand-built
Periodic table	153	326	4	2	Hand-built
Astronomy	41	39	6	6	Hand-built
Physics	43	46	6	5	Hand-built
Materials science	49	53	6	5	Hand-built
Wikidata Physics	252	814	12	7	Wikidata
Wikidata Materials	253	806	18	7	Wikidata

All statistics are computed programmatically from the KG builders via `analysis/reproducibility_table.py`. The dataset splitting protocol (Section 4) applies identically to all seven domains: 10% hidden, then 80/10/10 train/val/test on the remaining 90%.

B Ablation Details

The ablation study uses the linear algebra KG with $n_{\text{bootstrap}} = 200$ samples. For each ablation variant, one weight is set to zero and the remaining three are renormalised to sum to 1. Bootstrap 95% confidence intervals are computed via the percentile method on the mean Harmony score.

Weight sensitivity. We evaluate six weight configurations from the calibration gate grid ($\alpha \in \{0.3, 0.5, 0.7\}$, $\beta \in \{0.1, 0.3\}$, $\gamma = \delta = 0.25$). All configurations show Harmony > frequency baseline, with $\alpha = 0.5, \beta = 0.3$ yielding the highest mean Harmony score. This suggests that a moderate compressibility weight combined with non-trivial coherence weight best captures the structure of our curated KGs.

C Proposal Validation Rules

The deterministic validator enforces three rules:

- Text length:** claim, justification, and falsification_condition must each be ≥ 10 characters. kg_domain must be ≥ 3 characters (controlled vocabulary, not free text).
- Type-specific fields:** ADD_EDGE requires source_entity, target_entity, and edge_type; ADD_ENTITY requires entity_id and entity_type; REMOVE_EDGE requires source_entity, target_entity, and edge_type; REMOVE_ENTITY requires entity_id.
- Edge type validity:** edge_type must be one of the seven valid EdgeType names.

D Full Proposal Examples

Below are three complete proposal records from the astronomy archive, showing all fields including justification and falsification conditions.

Proposal 1: Stellar nucleosynthesis → heavy element abundance.

- **Type:** ADD_EDGE
- **Source:** stellar_nucleosynthesis
- **Target:** heavy_element_abundance
- **Edge type:** explains
- **Claim:** “Stellar nucleosynthesis explains the observed abundance pattern of heavy elements in planetary nebulae.”
- **Justification:** “The s-process and r-process nucleosynthesis pathways in AGB stars and supernovae produce characteristic abundance patterns that match spectroscopic observations of planetary nebulae.”
- **Falsification:** “Discovery of heavy element abundance patterns in planetary nebulae inconsistent with any known nucleosynthesis pathway would falsify this claim.”

Proposal 2: Mass–luminosity relation derivation.

- **Type:** ADD_EDGE
- **Source:** hydrostatic_equilibrium
- **Target:** mass_luminosity_relation
- **Edge type:** derives
- **Claim:** “The mass–luminosity relation derives from hydrostatic equilibrium in main sequence stars.”
- **Justification:** “Balancing gravitational pressure against radiation pressure in the stellar core, combined with opacity-dependent energy transport, yields $L \propto M^{3.5}$ for main sequence stars.”
- **Falsification:** “A main sequence star population where luminosity is uncorrelated with mass would disprove this derivation.”

Proposal 3: Magnetar as new entity.

- **Type:** ADD_ENTITY
- **Entity ID:** magnetar
- **Entity type:** celestial_object
- **Claim:** “Magnetar generalises the neutron star category with extreme magnetic field properties ($B > 10^{14}$ G).”
- **Justification:** “Magnetars are observationally distinct from ordinary neutron stars due to their ultra-strong magnetic fields, which power soft gamma repeaters and anomalous X-ray pulsars.”
- **Falsification:** “Evidence that magnetar-attributed emissions originate from non-magnetic mechanisms would undermine this classification.”

E LLM Prompt Templates

We include the exact prompt templates used for proposal generation. Both modes share a common preamble with KG statistics, strategy instruction, top proposals, and recent failures.

Free mode (default). The free-mode prompt shows a sample of up to 20 entity IDs from the KG to ground the LLM without over-constraining it:

```
You are a theory-discovery agent for knowledge graph research.
Knowledge Graph: domain='{domain}', entities={N}, edges={M}
Strategy: {REFINEMENT|COMBINATION|NOVEL} -- {strategy description}
Top proposals so far: {top 3 proposals or "None yet"}
Recent validation failures: {up to 5 failure messages or "None"}
```

```

347     EXAMPLE ENTITY IDs from this KG (showing K of N): {entity_1},
348     {entity_2}, ...
349     VALID EDGE TYPES: depends_on, derives, equivalent_to, maps_to,
350     explains, contradicts, generalizes
351     IMPORTANT: source_entity and target_entity MUST be exact entity IDs
352     from this KG.
353
354     Return ONLY a JSON object (no extra text) with fields: id,
355     proposal_type, claim, justification, falsification_condition,
356     kg_domain, source_entity, target_entity, edge_type, entity_id,
     entity_type

```

357 **Constrained mode (stagnation recovery).** When an island stagnates ($S = 5$ generations without
358 valid proposals), the prompt switches to constrained mode, which enumerates *all* valid entity IDs
359 and edge type names explicitly:

```

360     ... [same preamble] ...
361     VALID ENTITY IDs (use EXACTLY as written): {all entity IDs}
362     VALID EDGE TYPES (use EXACTLY as written): depends_on, derives,
363     equivalent_to, maps_to, explains, contradicts, generalizes

```

364 F Proposal Failure Rate Statistics

365 Figure 2 shows the valid proposal rate converging to ≥ 0.50 by generation 10 across all discovery
366 domains. The initial failure rate (generations 1–3) is typically 60–80%, dominated by entity ground-
367 ing errors (referencing entities not in the KG). The entity sample in free-mode prompts (up to 20
368 entities) and the stagnation recovery mechanism (Section 3.4) together reduce failures to $< 30\%$ by
369 generation 10. Constrained-mode prompts achieve $\geq 95\%$ validity but produce less diverse propos-
370 als.

371 G Code and Data Availability

372 Code and data are withheld for double-blind review and will be released upon acceptance:

- 373 • **Code repository:** [REDACTED FOR BLIND REVIEW] — full implementation of the Har-
374 mony metric, island-model search loop, MAP-Elites archive, and KG dataset builders.
- 375 • **Data archive:** [REDACTED FOR BLIND REVIEW] — all KG datasets, per-domain check-
376 points, generated proposals, and analysis artifacts.

377 H Statistical Methods

378 All pairwise comparisons between methods use the following statistical tests, computed over 10
379 random seeds per domain:

380 **Paired bootstrap CI.** Given paired score vectors $\mathbf{a} = (a_1, \dots, a_{10})$ and $\mathbf{b} = (b_1, \dots, b_{10})$, we
381 resample $B = 2000$ bootstrap differences $d^{(b)} = \frac{1}{10} \sum_i (a_{\pi_i} - b_{\pi_i})$ (sampling with replacement)
382 and report the 95% percentile CI $[d_{0.025}, d_{0.975}]$. The two-sided p -value is $p = 2 \cdot \min(\hat{P}(d^{(b)} \leq$
383 $0), \hat{P}(d^{(b)} \geq 0))$.

384 **Cliff’s delta.** A non-parametric effect size: $\delta_C = \frac{1}{n^2} \sum_{i,j} \text{sign}(a_i - b_j) \in [-1, 1]$. We interpret
385 $|\delta_C| < 0.147$ as negligible, < 0.33 as small, < 0.474 as medium, and ≥ 0.474 as large [3].

386 **Permutation test.** We randomly swap labels between methods 10,000 times to compute a Monte
387 Carlo estimate of the p -value under the null hypothesis of no difference.

I Frequency Baseline Analysis

For a KG with edge-type frequency vector $\mathbf{p} = (p_1, \dots, p_7)$, the frequency heuristic predicts the most common edge type for every unobserved entity pair. The expected Hits@ K of this strategy satisfies:

$$\text{Hits@}K \geq p_{\max} = \max_i p_i,$$

since for any test edge with the dominant type, the frequency prediction ranks the correct type first. In our domains, $p_{\max}$ ranges from 0.22 (Wikidata Materials, highest entropy) to 0.38 (Physics, lowest entropy), providing a lower bound on frequency Hits@10.

The Shannon entropy $H(\mathbf{p}) = -\sum_i p_i \log_2 p_i$ quantifies edge-type diversity; lower entropy indicates more peaked distributions where frequency is more informative. A regression of frequency Hits@10 on $H(\mathbf{p})$ across our seven domains shows a negative correlation, confirming the information-theoretic explanation.

J Wikidata Ablation

Per-component ablation on the Wikidata Physics and Wikidata Materials domains uses $n_{\text{bootstrap}} = 200$ resamples. On the denser Wikidata KGs, coherence shows a larger contribution than on the sparser hand-curated domains, consistent with the higher triangle density in larger graphs. Compressibility and symmetry contributions are broadly consistent across all domains. Generativity remains the single largest contributor, though the Harmony-3 analysis (Section 5) shows that the structural components alone provide useful signal.

K Rebuttal-Oriented Supplementary Tables

This appendix provides compact evidence tables intended to answer common review questions about stability, uncertainty, and runtime behavior. All tables are generated directly from analysis artifacts to avoid manual transcription errors.

K.1 MLX Batch Runtime Summary

Table 3: Wall-clock runtime (seconds) for MLX batch runs (mlx-community/Qwen3.5-35B-A3B-4bit via mlx_lm, Apple M2 Pro, 10 seeds per domain). Times include LLM inference, Harmony scoring, and checkpointing.

Domain	Mean (s)	Std (s)	Seeds	Mean (min)
Astronomy	1657.1	366.4	10/10	27.6
Physics	1375.8	215.2	10/10	22.9
Materials	1716.2	399.6	10/10	28.6
Wikidata Physics	1657.8	275.5	10/10	27.6
Wikidata Materials	2038.4	568.4	10/10	34.0

K.2 Factor Decomposition Detail

K.3 Statistical Tests Summary

L LLM Judge Evaluation Prompt

The following prompt (“detailed” variant from scripts/llm_judge_rubric.py) was used to score each proposal on four dimensions (1–5 scale). Proposal fields are sanitized (truncated and XML-stripped) before insertion. The paper reports scores generated via Gemini 2.5 Pro (Gemini CLI); llm_judge_rubric.py defaults to Claude for programmatic re-runs.

You are an expert scientific knowledge evaluator assessing theory proposals for

Table 4: Factor decomposition: best harmony gain per configuration and domain. LLM-only bypasses Harmony scoring (gain = 1.0 by construction, not comparable). Full = LLM proposer + Harmony scoring + MAP-Elites archive.

Domain	Configuration	Best Gain	Gens	Archive
Astronomy	LLM-only	—	11	1
	No-QD	0.0518	11	1
	Harmony-only	0.0539	20	5
	Full (LLM+Harmony+MAP-Elites)	0.0000	20	1
Physics	LLM-only	—	11	1
	No-QD	0.0259	12	1
	Harmony-only	0.0319	20	4
	Full (LLM+Harmony+MAP-Elites)	0.0278	13	3
Materials	LLM-only	—	11	1
	No-QD	0.0260	11	1
	Harmony-only	0.0291	20	5
	Full (LLM+Harmony+MAP-Elites)	0.0010	18	1
Wikidata Physics	LLM-only	—	11	1
	No-QD	0.0000	10	1
	Harmony-only	0.0021	20	4
	Full (LLM+Harmony+MAP-Elites)	0.0031	11	2
Wikidata Materials	LLM-only	—	11	1
	No-QD	0.0172	14	1
	Harmony-only	0.0186	20	5
	Full (LLM+Harmony+MAP-Elites)	0.0154	20	3

Table 5: Paired bootstrap CI (2000 resamples) and Cliff’s δ effect size for Harmony vs. DistMult-alone Hits@10. None reach $p < 0.05$. $|\delta_C| < 0.147$ = negligible; all observed $|\delta_C| \leq 0.12$.

Domain	$\Delta\text{Hits@10}$	95% CI	p	δ_C
Astronomy	+0.000	[0.000, 0.000]	1.000	+0.00
Physics	−0.044	[−0.200, 0.122]	0.620	−0.06
Materials	+0.020	[−0.050, 0.120]	0.718	+0.06
Wikidata Physics	+0.002	[−0.019, 0.023]	0.886	−0.06
Wikidata Materials	+0.012	[−0.019, 0.039]	0.378	+0.12

420 knowledge graphs. Evaluate the following proposal from the {domain}
421 domain.

422

423 <proposal>
424 <claim>{claim}</claim>
425 <justification>{justification}</justification>
426 <falsification_condition>{falsification}</falsification_condition>
427 <kg_mutation>{edge_type}({source} → {target})</kg_mutation>
428 </proposal>

429

430 RUBRIC:

431 Plausibility (1-5): Is the claim scientifically plausible given
432 known domain knowledge?
433 1=implausible/contradicts established knowledge, 3=uncertain/speculative,
434 5=well-grounded

435 Novelty (1-5): Does the claim add genuinely new information beyond
436 the existing KG?
437 1=trivially obvious/redundant, 3=moderately novel, 5=genuinely
438 surprising and new

439 Falsifiability (1-5): Is the falsification condition concrete and
440 experimentally testable?
441 1=unfalsifiable/vague, 3=partially testable, 5=precisely

operationalised
 Clarity (1-5): Is the claim stated clearly and specifically enough
 to be actionable?
 1=vague/ambiguous, 3=partially clear, 5=precise and unambiguous

Respond with ONLY a valid JSON object in this exact format:
 {"plausibility": <1-5>, "novelty": <1-5>, "falsifiability": <1-5>,
 "clarity": <1-5>}

No explanation, no markdown, no other text -- just the JSON object.

Evaluation script: scripts/llm_judge_rubric.py. Results stored in
 data/results/llm_rubric_scores.json. Top-5 proposals per domain were
 ranked by MAP-Elites archive fitness from the MLX batch run (10 seeds,
 mlx-community/Qwen3.5-35B-A3B-4bit).

M Hyperparameter Settings

Table 6 lists all hyperparameters used in the experiments.

Table 6: Hyperparameter settings.		
Component	Parameter	Value
Harmony metric	α (compressibility)	0.25
	β (coherence)	0.25
	γ (symmetry)	0.25
	δ (generativity)	0.25
	ε (frequency)	0.0
DistMult	Embedding dimension	50
	Training epochs	100
	Margin	1.0
	Learning rate	0.01
	Negative samples	5
	Mask ratio	0.20
Search loop	Islands	4
	Population per island	5
	Generations	20
	Migration interval	10
	Temperatures	{0.3, 0.3, 0.8, 1.2}
Stagnation	Trigger generations (S)	5
	Recovery generations (R)	3
MAP-Elites	Grid size	5×5
	Descriptors	simplicity, Harmony gain
Value function	λ (cost penalty)	0.1

Compute resources. All experiments were run on a single Apple M2 Pro (10-core CPU, 32 GB RAM, no GPU). Main experiments use gpt-oss:20b via Ollama (~ 10 min per domain, 20 generations). Factor decomposition baselines (LLM-only, No-QD) use mlx-community/Qwen3.5-35B-A3B-4bit via mlx_lm at ~ 46 tok/s, completing all three configurations across five domains in ~ 2.5 hours. The total compute for the seven reported domains is under 2 CPU-hours for the main experiments; factor decomposition adds ~ 2.5 CPU-hours. KGE baseline training (TransE, RotatE, ComplEx across 5 domains $\times$ 10 seeds) adds approximately 30 minutes. Statistical tests (bootstrap CI, permutation tests) add < 10 minutes.

N Declaration of LLM Usage

LLMs are core methodological components of this work: as the main proposer that generates KG mutations, as the proposer used in the factor decomposition study, and as a rubric judge for automated quality assessment. Table 7 consolidates the model identifiers, backends, and sampling parameters; verbatim prompt templates appear in Appendices E (proposer) and L (judge).

Table 7: LLM components, model identifiers, backends, and sampling parameters.

Component	Model ID	Backend	Sampling
Proposer (main)	<code>gpt-oss:20b</code>	Ollama	$\text{temp} \in \{0.3, 0.3, 0.8, 1.2\}$ per island
Proposer (factor dec.)	<code>mlx-community/Qwen3.5-35B-A3B-4bit</code>	<code>mlx_lm</code>	$\text{temp } 0.7$, $\text{top-}k \ 50$, $\text{top-}p \ 0.9$, $\text{max_tok } 1024$
Judge (rubric)	Gemini 2.5 Pro	Gemini CLI	defaults

The 10 random seeds used across all multi-seed runs are:

$\{42, 123, 456, 789, 1024, 1337, 2048, 3141, 4096, 5000\}$.

MLX backend defaults ($\text{max_tokens}=1024$, $\text{temperature}=0.7$, $\text{top_k}=50$, $\text{top_p}=0.9$) are set in the supplementary code archive; the Ollama backend passes only temperature (per-island) and uses Ollama defaults for the remaining sampler parameters (top_k , top_p , repeat_penalty). Self-correction repair calls fix temperature at 0.3.

Software environment. Reproductions depend on the runtime versions of the inference engines. Main runs used Ollama ≥ 0.3 with the `gpt-oss:20b` model (SHA-256 manifest pinned in the supplementary archive); factor decomposition used `mlx-lm` ≥ 0.20 with the linked HuggingFace model card. Gemini 2.5 Pro is invoked via the public Gemini CLI without a fixed checkpoint; the rubric output is reported only as a secondary diagnostic.

O Edge Type Rationale

The seven relation types are derived from a morphism-first principle: we surveyed the core semantic roles needed to express scientific relationships across five domains (linear algebra, chemistry, astronomy, physics, materials science) and identified a minimal set that covers dependency (`depends_on`), derivation (`derives`), equivalence (`equivalent_to`), correspondence (`maps_to`), causal/explanatory links (`explains`), contradiction (`contradicts`), and taxonomic hierarchy (`generalizes`). These seven types are inspired by morphism classes in category theory, and we found that scientific relations across our five evaluation domains map naturally to one of these types. The fixed vocabulary enables cross-domain comparisons while remaining expressive enough to capture the core semantic relations in scientific knowledge.

P Formal Properties

The Harmony metric satisfies three properties that make it suitable as a discovery prior:

1. **Boundedness:** $\mathcal{H}(G) \in [0, 1]$ for any KG G , since each component is bounded in $[0, 1]$ and weights are normalised to sum to 1.
2. **Decomposability:** each component (Compress, Cohere, Symm, Gener) is independently computable from the graph structure, enabling parallel evaluation and interpretable ablation.
3. **Directional monotonicity** (empirical observation): each component *tends to* respond predictably to edge addition—compressibility generally decreases (more cross-edges reduce spanning fraction), coherence increases when the new edge closes a type-consistent triangle, symmetry increases when the edge balances entity-type distributions, and generativity increases when the edge adds learnable relational signal. The Harmony score thus captures the *net* structural effect of a mutation across these competing pressures. We note




Figure 3: MAP-Elites archive fitness heatmaps. Each cell shows the fitness of the elite proposal at that (simplicity, gain) bin. Empty cells (white) indicate unexplored regions of the behavioural space.

that these are empirical tendencies, not formal guarantees; edge placement can produce non-monotonic effects in individual components.

Philosophical grounding. The four components correspond to established principles of theory quality: compressibility instantiates Occam’s razor via minimum description length (MDL); coherence enforces logical consistency across relational paths; symmetry operationalises an intuition analogous to Noether’s theorem—that good theories exhibit invariance across structurally equivalent entities; and generativity captures predictive validity—the hallmark of a useful scientific theory.

Q Calibration Gate Results

The calibration gate passed on both domains. On the linear algebra KG, the Harmony score exceeds the frequency baseline by 31% (bootstrap 95% CI: [0.24, 0.38]). On the periodic table KG, the improvement is 65% (95% CI: [0.52, 0.78]). All six pre-registered weight configurations show consistent direction (Harmony > frequency), confirming that the metric’s advantage is robust to weight choices.

R Proposal Validity and Archive Coverage

Across all five discovery domains (astronomy, physics, materials, Wikidata Physics, and Wikidata Materials), the valid proposal rate reaches ≥ 0.50 by generation 10, satisfying the pre-registered gate condition. The MAP-Elites archive achieves 40–60% coverage of the 5×5 grid (10–15 of 25 cells occupied), indicating that the island-model search produces diverse proposals spanning multiple simplicity–gain trade-offs (Figure 3).

S Factor Decomposition Detail

To understand what each component contributes, we compare the full Harmony pipeline against three ablated configurations (Section 4): LLM-only (bypass Harmony scoring, `-accept-all-valid`), Harmony-only (random proposer, no LLM), and No-QD (greedy single-bin archive, `-greedy`). Table 8 reports the best harmony gain achieved by each configuration across 20 generations.

Three findings emerge. First, **Harmony-only outperforms No-QD on all five domains** despite using a random proposer instead of an LLM. On the four domains where No-QD achieves non-zero gain, Harmony-only yields 4–23% higher best gain; on Wikidata Physics, No-QD stagnates at zero gain while Harmony-only reaches 0.002. This demonstrates that the MAP-Elites quality-diversity archive is the primary driver of structural gain: by maintaining diverse proposals across the simplicity–gain grid, the archive avoids premature convergence (No-QD early-stops at generation 10–14; Harmony-only sustains improvement through generation 20).

Second, the **LLM-only** configuration (all valid proposals accepted, no Harmony scoring) saturates the archive trivially—every proposal receives a constant gain of 1.0 by construction. The lack of selection pressure means the archive cannot distinguish high-quality from low-quality proposals, confirming that Harmony scoring provides essential signal for proposal ranking.

Third, the **Full pipeline** (LLM + Harmony + MAP-Elites) achieves gain comparable to or slightly below Harmony-only on most domains. On Wikidata Physics, Full (0.003) marginally exceeds




Figure 4: Hits@10 comparison across discovery domains. TransE achieves the highest Hits@10 on three domains; the frequency heuristic leads on Astronomy and Wikidata Materials. Harmony-guided proposals (green) show marginal improvements over the DistMult-alone embedding baseline in three of five domains (no differences are statistically significant at $\alpha = 0.05$).

Harmony-only (0.002); on Physics, Full (0.028) is close to Harmony-only (0.032). On Astronomy, the single Full run did not converge to positive gain within the generation budget, while Harmony-only reached 0.054; on Materials, Full achieved only negligible gain (0.001) versus Harmony-only’s 0.029. This indicates the LLM proposer’s marginal contribution to raw Harmony gain is small relative to the MAP-Elites archive. The LLM’s primary value is *semantic*—producing proposals with structured justifications and falsification conditions—rather than maximising structural gain directly.

Together, these results show that Harmony scoring and MAP-Elites diversity are the dominant drivers of proposal quality; the LLM proposer adds semantic structure that the random proposer cannot provide.

T Backtesting Detail

Table 9 reports backtesting results under two criteria against 10%-held-out hidden edges (up to 50 proposals pooled across 10 seeds per domain from the MLX batch run). *Exact match*: all three of source, target, and edge type must match a hidden triple. *Soft match*: at least one of (source, edge_type) or (edge_type, target) matches a hidden edge, testing whether proposals target structurally relevant neighbourhoods.

Exact match precision and recall are zero across all domains, confirming that LLM-generated proposals do not reproduce held-out edges. This serves as a *negative-control diagnostic*: proposals behave as theory-level mutations rather than interpolations of existing triples.

Soft-match results reveal that proposals do reach structurally relevant neighbourhoods in some domains. On Wikidata Physics, 20% of top-10 proposals share a (source, edge_type) or

Table 8: Factor decomposition: best harmony gain per configuration. LLM-only bypasses scoring (gain trivially = 1.0, not comparable). No-QD uses a greedy single-bin archive; Harmony-only uses a random proposer with the full 5×5 MAP-Elites archive; Full uses the LLM proposer with Harmony scoring and the MAP-Elites archive. Archive size = occupied MAP-Elites cells at termination.

Domain	Config	Best Gain	Gens	Archive
Astronomy	LLM-only	—	11	1
	No-QD	0.052	11	1
	Harmony-only	0.054	20	5
	Full	0.000	20	1
Physics	LLM-only	—	11	1
	No-QD	0.026	12	1
	Harmony-only	0.032	20	4
	Full	0.028	13	3
Materials	LLM-only	—	11	1
	No-QD	0.026	11	1
	Harmony-only	0.029	20	5
	Full	0.001	18	1
Wikidata Physics	LLM-only	—	11	1
	No-QD	0.000	10	1
	Harmony-only	0.002	20	4
	Full	0.003	11	2
Wikidata Materials	LLM-only	—	11	1
	No-QD	0.017	14	1
	Harmony-only	0.019	20	5
	Full	0.015	20	3

Table 9: Backtesting: exact and soft (typed-neighbourhood) match against 10% held-out hidden edges. Exact P/R require all three fields to match; Soft P/R require sharing (`source`, `edge_type`) or (`edge_type`, `target`) with a hidden triple. Soft P uses @10 (precision of the top-10 proposals); Soft R uses @20 (recall over a wider candidate set to better cover hidden edges). Evaluated at the top- N proposals pooled across 10 seeds.

Domain	Exact P@20	Exact R@20	Soft P@10	Soft R@20
Astronomy	0.00	0.00	0.00	0.25
Physics	0.00	0.00	0.10	0.20
Materials	0.00	0.00	0.00	0.00
Wikidata Physics	0.00	0.00	0.20	0.02
Wikidata Materials	0.00	0.00	0.10	0.05

(`edge_type`, `target`) pair with a held-out edge. On Astronomy, 25% of held-out edges are “soft-covered” by top-20 proposals. Materials shows zero soft match, suggesting proposals in that domain are more exploratory and structurally distant from the held-out set. These results are consistent with the interpretation that Harmony-guided proposals tend toward structurally adjacent knowledge rather than being purely random, even when they do not reproduce exact triples. Notably, the zero exact-match result holds even for the denser Wikidata domains (252–253 entities, 806–814 edges); this consistency across both sparse hand-curated and medium-scale KGs strengthens the negative-control interpretation.

U Regime Characterisation

Rather than claiming universal superiority, we characterise the *regimes* where Harmony adds value beyond frequency. Plotting KG density (edges per entity) against the Harmony–frequency Hits@10 gap across all seven domains reveals a trend: Harmony’s complementary value increases with graph density and edge-type diversity. On sparse, low-entropy KGs (e.g. hand-curated Physics), frequency captures the dominant pattern trivially; on denser, higher-entropy KGs (e.g. Wikidata Materials), the structural signals in the Harmony metric identify proposals that frequency misses.

Table 10: Automated proposal quality scores: mean per domain across top-5 proposals (scored by Gemini 2.5 Pro; 1–5 scale). Plausibility = scientific groundedness; Novelty = information gain beyond common knowledge; Falsifiability = testability of the falsification condition; Clarity = specificity of the claim.

Domain	Plausibility	Novelty	Falsifiability	Clarity
Astronomy	4.0	2.8	5.0	5.0
Physics	2.6	3.0	3.8	3.6
Materials	3.5	2.5	5.0	5.0
Wikidata Physics	3.4	3.2	4.0	4.8
Wikidata Materials	4.4	2.6	4.6	4.8
Overall (mean)	3.6	2.8	4.5	4.6

V Automated Quality Assessment

Table 10 reports automated quality scores for top-5 proposals per domain, scored on four dimensions (1–5 scale: 1=poor, 5=excellent) using Gemini 2.5 Pro as an automated evaluator. This follows the LLM-as-judge paradigm [15]: automated scoring is scalable, reproducible, and consistent within a single model, though it does not replace human expert evaluation with inter-annotator agreement.

Across domains, falsifiability (4.5) and clarity (4.6) are consistently high, confirming that the structured proposal schema (Section 3.3) reliably produces testable, specific claims. Plausibility is highest for Wikidata Materials (4.4) and Astronomy (4.0), where domain knowledge is more concrete; Physics scores lower (2.6) because its top proposals include speculative claims linking magnetic moment topology to gravitational regimes — a domain where the LLM explores theoretical edges of known physics. Novelty is moderate across domains (2.5–3.2), reflecting the expected tension between semantic grounding (proposals must reference entities in the KG) and genuine theoretical novelty.

Disclosure: Scores are from a single LLM judge (Gemini 2.5 Pro) with no inter-annotator agreement measurement. Human expert evaluation with domain specialists and measured Cohen’s κ is deferred to future work. The rubric prompt is included in Appendix L for reproducibility.

W Extended Discussion

Sparse KG challenges. Our curated KGs range from 41–253 entities and 39–814 edges across all seven domains (Table 2). The smaller hand-curated domains (41–49 entities) represent early-stage scientific KG construction where sparsity limits the generativity component: DistMult requires ≥ 10 training edges to produce meaningful predictions, and the 20% masking protocol leaves few test edges for evaluation on these domains. The Wikidata domains (252–253 entities, 806–814 edges) confirm that the framework scales to medium-sized KGs.

Proposal quality vs. validity rate. The stagnation recovery mechanism (constrained prompting after $S = 5$ generations without valid proposals) effectively maintains a validity rate ≥ 0.50 across domains. However, constrained proposals tend to cluster in low-novelty regions of the MAP-Elites grid. A promising direction is adaptive constraint relaxation, where the degree of structural constraint is modulated by archive coverage rather than a binary switch.

Symmetry redesign rationale. In the initial design, symmetry measured inter-type behavioural uniformity via pairwise JS divergence across entity types. Reviewer feedback correctly identified that this penalises natural functional specialisation: in an astronomy KG, stars and planets *should* exhibit different edge-type distributions. The revised design (Eq. 4) measures *intra-type* behavioural consistency—whether entities of the same type use edge types consistently—which is a more meaningful structural signal. Entities with no outgoing edges are excluded, and single-entity types receive maximal consistency by convention.

Contradicts validity. Contradicts edges need not represent noise—in scientific discourse, competing hypotheses are valuable and their explicit representation is a feature, not a defect. Our

Table 11: Representative proposals from the astronomy MAP-Elites archive.

Type	Edge type	Claim
ADD_EDGE	explains	“Stellar nucleosynthesis explains the observed abundance pattern of heavy elements in planetary nebulae.”
ADD_EDGE	derives	“The mass–luminosity relation derives from hydrostatic equilibrium in main sequence stars.”
ADD_ENTITY	—	“Magnetar (entity type: celestial_object) generalises the neutron star category with extreme magnetic field properties.”

coherence penalty targets only *dense* contradiction (high contradicts-to-edge ratio), which signals structural noise; sparse contradiction is tolerated. Future work includes domain-adaptive weighting, where component weights are learned per domain via held-out validation performance.

Harmony–downstream disconnect. Per-proposal Spearman rank correlations between component deltas and $\Delta\text{Hits}@10$ would require logging per-proposal component scores alongside evaluation results. Our current event logging captures only generation-level summaries, so this analysis is deferred to future work. We view this as an important finding that motivates future metric refinement rather than a fatal flaw: the Harmony metric captures structural desiderata that are valuable for theory assessment beyond link prediction alone.

Scalability. The Harmony framework’s computational cost is dominated by DistMult training ($O(|E| \cdot d \cdot \text{epochs})$) and LLM inference ($O(T_{\max} \cdot 4)$ calls for 4 islands). The three graph-structural components (compressibility, coherence, symmetry) are $O(|V| + |E|)$ each. For our hand-curated KGs (41–49 entities, 39–53 edges), total wall time is ~ 10 minutes per domain on a single CPU. The Wikidata domains (252–253 entities, 806–814 edges) complete in ~ 27 minutes per seed (Table 3). Scaling to medium-size KGs ($\sim 1,000$ entities) increases DistMult training time linearly with $|E|$ but does not change the LLM call count. However, constrained prompting (Section 3.4) enumerates all valid entity IDs in the prompt, so prompt token volume grows with $|V|$; mitigations include chunked entity sampling or index-based lookups. The framework is practical without GPU hardware for KGs up to a few thousand entities.

X Qualitative Examples (Extended)

Table 11 shows representative proposals from the astronomy domain, illustrating the diversity of claims and mutation types.

References

- [1] Jinheon Baek, Alham Fikri Aji, and Amir Saffari. Knowledge-augmented language model prompting for zero-shot knowledge graph question answering. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 8696–8704. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.findings-emnlp.580.
- [2] Antoine Bordes, Nicolas Usunier, Alberto Garcia-Duran, Jason Weston, and Oksana Yakhnenko. Translating embeddings for modeling multi-relational data. In *Advances in Neural Information Processing Systems*, volume 26. Curran Associates, Inc., 2013. URL <https://proceedings.neurips.cc/paper/2013/hash/1cecc7a77928ca8133fa24680a88d2f9-Abstract.html>.
- [3] Norman Cliff. Dominance statistics: Ordinal analyses to answer ordinal questions. *Psychological Bulletin*, 114(3):494–509, 1993. doi: 10.1037/0033-2909.114.3.494.
- [4] Miles Cranmer. Interpretable machine learning for science with PySR and SymbolicRegression.jl, 2023. URL <https://arxiv.org/abs/2305.01582>.

- [5] Shaoxiong Ji, Shirui Pan, Erik Cambria, Pekka Marttinen, and Philip S. Yu. A survey on knowledge graphs: Representation, acquisition, and applications. *IEEE Transactions on Neural Networks and Learning Systems*, 33(2):494–514, 2022. doi: 10.1109/TNNLS.2021.3070843.
- [6] Jinhao Jiang, Kun Zhou, Zican Dong, Keming Ye, Wayne Xin Zhao, and Ji-Rong Wen. Struct-GPT: A general framework for large language model to reason over structured data. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9237–9251. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.574.
- [7] John R. Koza. *Genetic Programming: On the Programming of Computers by Means of Natural Selection*. MIT Press, Cambridge, MA, 1992. ISBN 978-0-262-11170-6.
- [8] Nour Makke and Sanjay Chawla. Interpretable scientific discovery with symbolic regression: A review, 2024. URL <https://arxiv.org/abs/2211.10873>.
- [9] Jean-Baptiste Mouret and Jeff Clune. Illuminating search spaces by mapping elites, 2015. URL <https://arxiv.org/abs/1504.04909>.
- [10] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M. Pawan Kumar, Emilien Dupont, Francisco J. R. Ruiz, Jordan S. Ellenberg, Pengming Wang, Omar Fawzi, Pushmeet Kohli, and Alhussein Fawzi. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024. ISSN 1476-4687. doi: 10.1038/s41586-023-06924-6.
- [11] Jiashuo Sun, Chengjin Xu, Luminyuan Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Heung-Yeung Shum, and Jian Guo. Think-on-graph: Deep and responsible reasoning of large language model on knowledge graph. In *Proceedings of the International Conference on Learning Representations (ICLR)*. arXiv, 2024. doi: 10.48550/arXiv.2307.07697. URL <https://arxiv.org/abs/2307.07697>.
- [12] Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019. doi: 10.48550/arXiv.1902.10197.
- [13] Darrell Whitley, Soraya Rana, and Robert B. Heckendorn. *Island model genetic algorithms and linearly separable problems*, pages 109–125. Springer Berlin Heidelberg, 1997. ISBN 9783540695783. doi: 10.1007/bfb0027170.
- [14] Bishan Yang, Wen-tau Yih, Xiaodong He, Jianfeng Gao, and Li Deng. Embedding entities and relations for learning and inference in knowledge bases. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2015. doi: 10.48550/arXiv.1412.6575.
- [15] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36. Curran Associates, Inc., 2023. doi: 10.48550/arXiv.2306.05685. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41f33-Abstract-Datasets_and_Benchmarks.html.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state three contributions: (1) the Harmony metric (Section 3), (2) the island-model search loop (Section 3), and (3) LLM-guided proposal generation (Section 3). Section 5 validates each with quantitative results.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 contains a dedicated “Limitations” paragraph addressing concrete constraints including coarse relation vocabulary, metric-domain mismatch risk, evaluation circularity caveats, and lack of statistical significance on most domains despite multi-seed evaluation.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper does not claim formal theorems. The Harmony metric (Eqs. 1–6) is defined compositionally; all component definitions and normalisation conventions are stated explicitly in Section 3.

Guidelines:

- The answer NA means that the paper does not include theoretical results.

- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: Appendix M lists all hyperparameters; Appendix A provides dataset statistics; Section 4 describes the experimental protocol; all 10 random seeds are listed explicitly.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [\[Yes\]](#)

Justification: Code is provided as a supplementary ZIP archive accompanying the submission; see Appendix G. The Zenodo DOI and de-anonymised public repository are deferred to camera-ready.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 describes the experimental setup including baselines, evaluation protocol, and dataset split ratios. Appendix M lists all hyperparameters. Appendix A provides dataset entity and edge count statistics.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 1 reports mean $\pm$ std across $n = 10$ seeds for all main results. Pairwise comparisons against DistMult-alone use paired bootstrap 95% CIs ($B = 2000$ resamples) and Cliff's δ for non-parametric effect size; the permutation-test protocol and full definitions are in Appendix H. Variability sources captured by the standard deviations: dataset split, model initialisation, edge masking.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix M reports the hardware (CPU-only, Apple M-series) and wall-clock time (approximately 10 minutes per domain for 20 generations). No GPU resources were used.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: No human subjects were involved. All knowledge graphs are curated from publicly available academic sources. No personally identifiable or scraped data is used.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 6 includes a "Broader impacts" paragraph discussing positive impacts (accelerating scientific theory discovery) and negative risks (LLM-generated claims may be plausible-sounding but factually incorrect, requiring expert validation before use in downstream scientific workflows).

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release pretrained models or scraped datasets. The released assets are small curated knowledge graphs and search-loop code, which pose no misuse risk.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: DistMult [14], TransE [2], RotatE [12], and ComplEx are cited. Core Python libraries (NumPy, scikit-learn, MLX, Ollama) are released under permissive licences (BSD-3-Clause, MIT, Apache-2.0). The five hand-curated KG datasets are original to this work; the two Wikidata-derived KGs (Wikidata Physics, Wikidata Materials) are subsets of the Wikidata knowledge graph, which is released under the Creative Commons CC0 1.0 licence (with attribution to Wikidata recommended).

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: Five curated KG datasets are documented in Appendix A with entity/edge counts, type vocabularies, and split ratios. The proposal schema is defined in Section 3 with validation rules in Appendix C.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: No crowdsourcing or research with human subjects was conducted.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

1015 Answer: [N/A]

1016 Justification: No human subjects research was conducted; IRB approval is not applicable.

1017 Guidelines:

- 1018 • The answer NA means that the paper does not involve crowdsourcing nor research
- 1019 with human subjects.
- 1020 • Depending on the country in which research is conducted, IRB approval (or equiva-
- 1021 lent) may be required for any human subjects research. If you obtained IRB approval,
- 1022 you should clearly state this in the paper.
- 1023 • We recognize that the procedures for this may vary significantly between institutions
- 1024 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
- 1025 guidelines for their institution.
- 1026 • For initial submissions, do not include any information that would break anonymity
- 1027 (if applicable), such as the institution conducting the review.

1028 **16. Declaration of LLM usage**

1029 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1030 non-standard component of the core methods in this research? Note that if the LLM is used

1031 only for writing, editing, or formatting purposes and does not impact the core methodology,

1032 scientific rigorousness, or originality of the research, declaration is not required.

1033 Answer: [Yes]

1034 Justification: LLM-based proposal generation is a core methodological component (Sec-

1035 tion 3). The proposer (gpt-oss:20b via Ollama, island temperatures {0.3, 0.3, 0.8, 1.2}),

1036 the factor-decomposition proposer (mlx-community/Qwen3.5-35B-A3B-4bit via

1037 mlx_lm, temp 0.7), and the rubric judge (Gemini 2.5 Pro) are consolidated in Appendix N

1038 (Table 7), with verbatim prompts in Appendices E and L.

1039 Guidelines:

- 1040 • The answer NA means that the core method development in this research does not
- 1041 involve LLMs as any important, original, or non-standard components.
- 1042 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>)
- 1043 for what should or should not be described.